


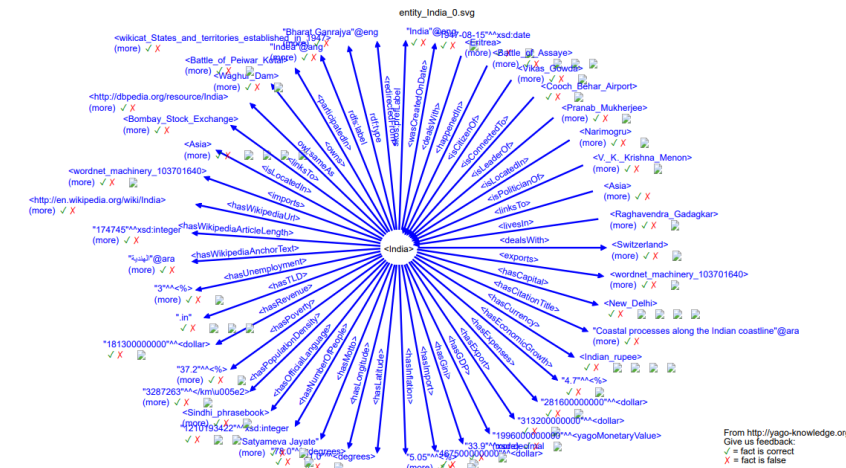


# Distributed Knowledge Graph Querying on Edge and Cloud

**DS 256 – Project Final Presentation**  
**Shriram R**

# Motivation

- Knowledge Graphs are large
- Graph querying problems are hard
- IoT – Heterogenous Structure
  - Different Compute, Storage & Network Capacity
- Needs lightweight and low latency solution





# Related Work

## ■ Graph Processing

- ▶ Pregel [2]
- ▶ Giraph [3]
- ▶ GraphX [4]
- ▶ Trinity [5]
- ▶ Quegel [13]

## ■ Graph Caching

- ▶ GC [6]

## ■ Graph Indexing

- ▶ C-Tree [7]
- ▶ Views [8]
- ▶ FERRARI [9]
- ▶ GraphS [10]

## ■ Graph Databases

- ▶ GoDB [11]
- ▶ TinkerGraph
- ▶ Neo4j
- ▶ Dgraph
- ▶ Titan


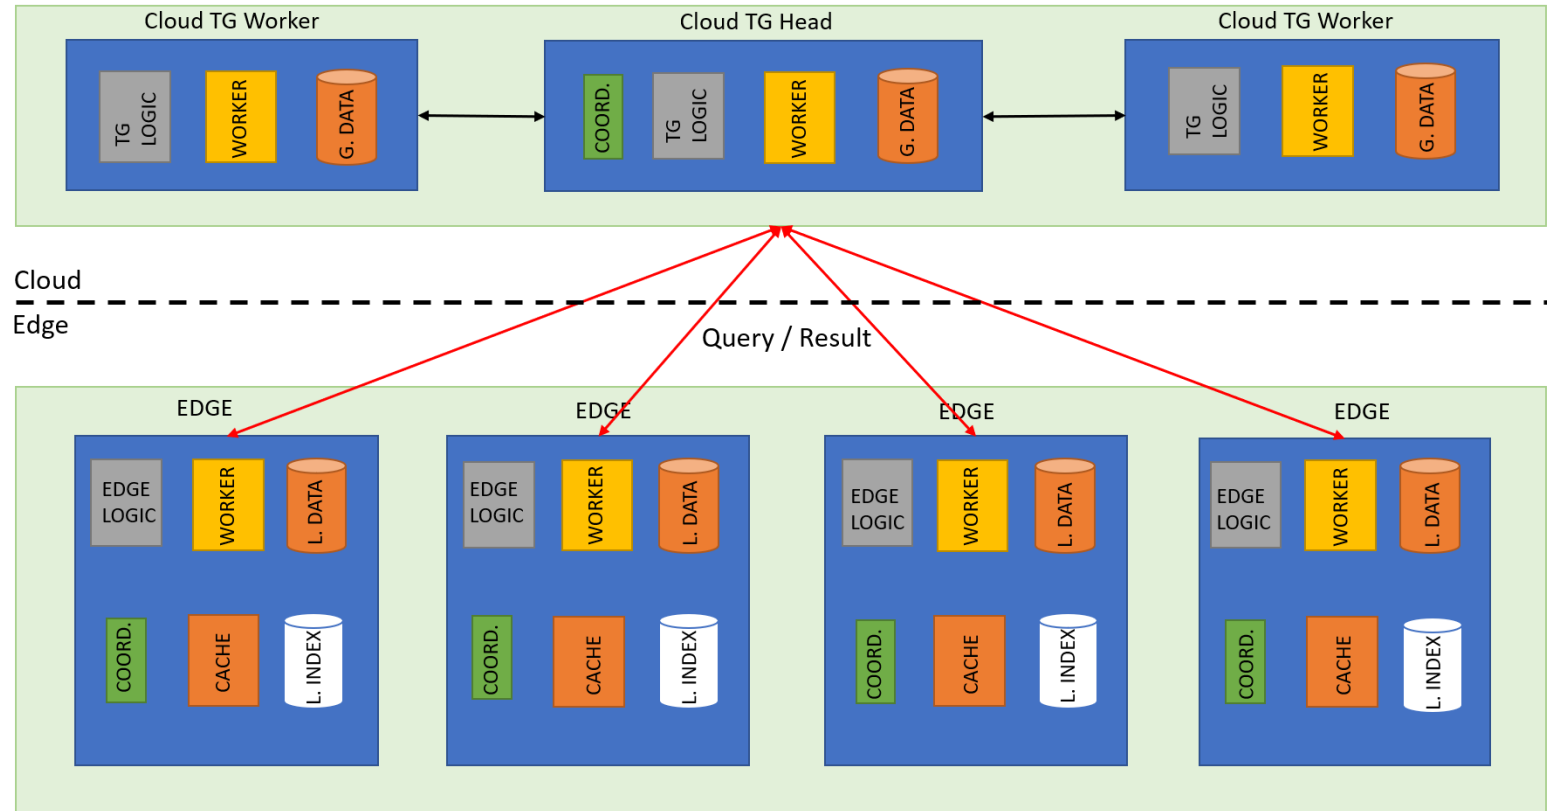


# Key Contributions

- Graph Caching
  - Cache Knowledge graph in Edge layer
- Query Partitioning
  - Partition input graph query into local and remote
- Result Generation
  - Combine results from local and remote server into a correct result
- Different Query Types
  - Vertex, Edge, Path search and reachability
- Result Caching
  - Cache results from remote layer with cache management




# System Architecture



# Graph Caching

- Associate an entity with the edge device E.g. <India>
- Store the **k-hop** subgraph around the entity as local graph
- Parameter **k** depends on storage and compute capacity





# Query Processing

## ■ Reachability

- ▶ Type 1 - Entirely contained within edge (local graph)
- ▶ Type 2 - Crosses between edge and remote (local + remote graph)
- ▶ Type 3 - Entirely contained within remote (remote graph)

## ■ Steps

1. Local Query is fired to check for Type 1 existence
2. If no result in previous step, remote queries are fired with source as Cut Vertices for Type 2
3. Local queries are fired with source unchanged and target as Cut Vertices
4. Remote query is also fired with original source for Type 3

# Query Processing

## ■ Path Query

- ▶ *N* vertex and *N-1* edge predicates
- ▶ *Null* in query is wildcard
- ▶ If path length  $< k$  (hops)
  - Fire local and remote query
- ▶ Else
  - Fire only remote query
- ▶ Union paths from local and remote query for the result

// Path Query Example

```
"type": "path_search",  
"filter": {  
  "vertices": ["Mahatma_Gandhi", null, "Subhash_Agarwal"],  
  "edges": ["isCitizenOf", null]  
}
```





# Result Caching

- Edge device caches query results from remote server
- Cache size in terms of bytes is parameterized
- Whenever a remote query is fired, cache is checked for hit
- New query is inserted with key as query and value as result
- FIFO Replacement policy
- Local queries are not cached




# Implementation

- **TinkerGraph** used for both Cloud and Edge
- Edge logic written in *Python*
- *GremlinPython* package used for communication
- *GraphML* format used to persist data


# Experiments

- Remote Graph database
  - 1 node in-memory spawned TinkerGraph in Rigel Head Node
  - 32 core AMD Opteron with 128 GB RAM
  - Centos 7
- Edge Layer
  - 4 core, 1GB RAM container
  - Ubuntu 18.04 LTS
  - Latency: 5ms & Bandwidth: 100 MBps
- Dataset
  - YAGO minimal knowledge graph
  - 14977 vertices, 18845 edges and 1400 total attributes
  - Local graph centred on <India> with 2 hops

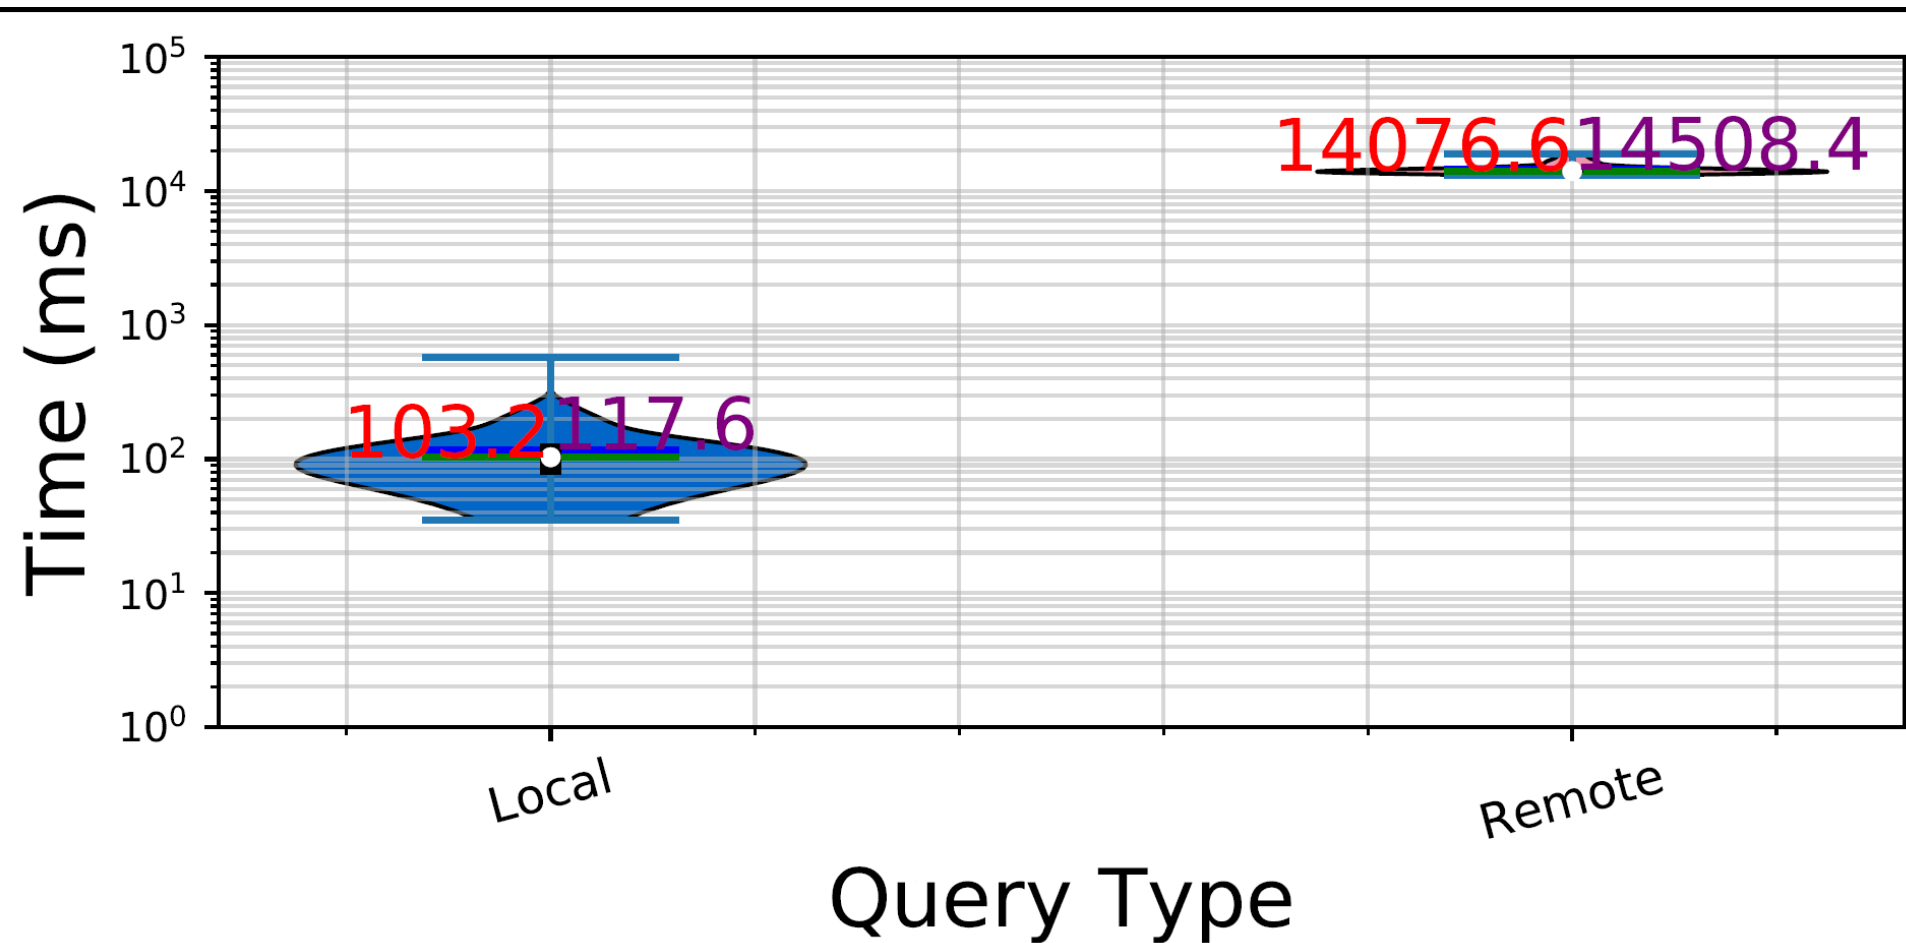
# Experiments – Vertex Search



# Experiments – Edge Search



# Experiments – Reachability





# Challenges & Future Work

## ■ Challenges

- ▶ Integrating *GoDB* with the project
- ▶ Setting up distributed environment for *TinkerGraph*
- ▶ Difficulty in setting up large cluster for experiments
- ▶ ETL for the dataset

## ■ Future Work

- ▶ Evaluate caching and all queries in large setup
- ▶ Index the graph on the edge layer
- ▶ Integrate with different datasets
- ▶ Implement machine learning algorithms using this system



# Original Proposal & Accomplishments

## ■ Midterm

- ▶ Graph Caching
- ▶ Query Partitioning
- ▶ Combining results
- ▶ Vertex, Edge search and Reachability
- ▶ Experiments

## ■ Final

- ▶ Caching of remote results
- ▶ Support for path queries
- ▶ Indexing
- ▶ Experiments

## ■ Stretch Goals

- ▶ Subgraph isomorphism queries
- ▶ Support for Graph updates